

Reinforcement Learning and Applications


Seung Hyun Oh

20240305

Outlines

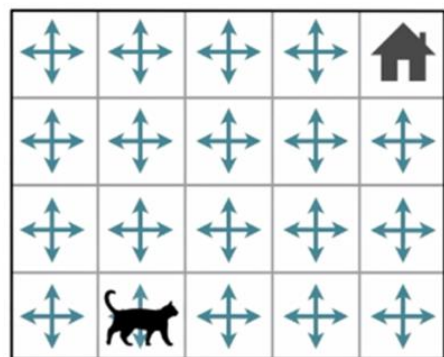
- What is Reinforcement Learning?
- Q-Learning
- DQN/DDPG/PPO
- Recent RL Algorithms(ex : MOPPO, TD3 ...)
- Applications (ex : Vessel) // Vessel Dynamics (Action, State), 적용되는지

Reinforcement Learning




- Find optimal action sequence for expected reward maximization
- The agent starts in a given state within its environment $s_0 \in S$ by gathering an initial observation $\omega_0 \in \Omega$
- At each time step t ,
The agent has to take an action $a_t \in A$

Q-learning



Behavior policy



Target policy

→ 평가하고 업데이트하고자 하는 policy

→ 실제로 행동해서 다음 state를 얻는 것

Q-value


$$Q(s, a) \leftarrow \underset{\text{Update}}{(1 - \alpha)} Q(s_t, a_t) + \underset{\text{Damping}}{\alpha} (R_t + \sigma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}))$$

$$= \int_{s_{t+1} a_{t+1}} (R_t + \sigma Q(s_{t+1}, a_{t+1})) \underset{\text{Target Policy}}{p(a_{t+1} | s_{t+1})} \underset{\text{R.V}}{p(s_{t+1} | s_t, a_t)} ds_{t+1}, a_{t+1}$$


$$\doteq Q_N = (1 - \alpha) Q_{N-1} + \alpha (R_t^N + \sigma Q(s_{t+1}^n, a_{t+1}^n))$$

DQN (Deep Q-learning Network)

Classic Q-learning



Deep Q-learning



- Q-function

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

탐험을 감가율

- Loss Function & Update (SGD)

$$L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot)} \left[(y_i - Q(s, a; \theta_i))^2 \right]$$

가중치 미래 보상 추정치 예측값



$$y_i = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) \mid s, a \right]$$

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right].$$

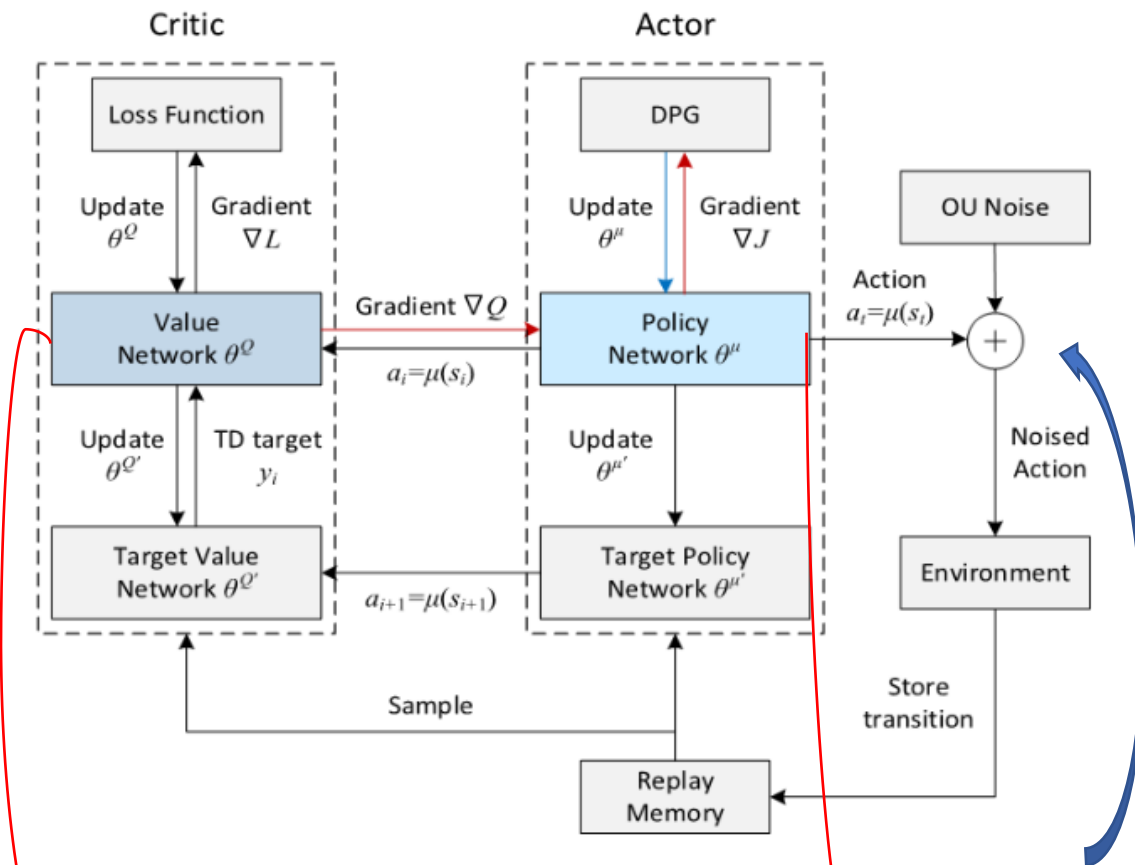
Algorithm 1 Deep Q-learning with Experience Replay

```

Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights
for episode = 1,  $M$  do
  Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$ 
  for  $t = 1, T$  do
    With probability  $\epsilon$  select a random action  $a_t$ 
    otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$ 
    Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
    Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
    Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$ 
    Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$ 
    Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$ 
    Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  according to equation 3
  end for
end for
  
```



DDPG(Deep Deterministic Policy Gradient)



$$L(\theta^Q) = \mathbb{E}_{s_t \sim \rho^\beta, a_t \sim \beta, r_t \sim E} [(Q(s_t, a_t | \theta^Q) - y_t)^2]$$


Loss Function



$$\begin{aligned} \nabla_{\theta^\mu} J &\approx \mathbb{E}_{s_t \sim \rho^\beta} [\nabla_{\theta^\mu} Q(s, a | \theta^Q) |_{s=s_t, a=\mu(s_t | \theta^\mu)}] \\ &= \mathbb{E}_{s_t \sim \rho^\beta} [\nabla_a Q(s, a | \theta^Q) |_{s=s_t, a=\mu(s_t)} \nabla_{\theta^\mu} \mu(s | \theta^\mu) |_{s=s_t}] \end{aligned}$$

Policy Gradient

- Q-function



$$Q^\pi(s_t, a_t) = \mathbb{E}_{r_t, s_{t+1} \sim E} [r(s_t, a_t) + \gamma \mathbb{E}_{a_{t+1} \sim \pi} [Q^\pi(s_{t+1}, a_{t+1})]]$$

$$Q^\mu(s_t, a_t) = \mathbb{E}_{r_t, s_{t+1} \sim E} [r(s_t, a_t) + \gamma Q^\mu(s_{t+1}, \mu(s_{t+1}))]$$

DDPG(Deep Deterministic Policy Gradient)

Algorithm 1 Deep Q-learning with Experience Replay

```

Initialize replay memory  $\mathcal{D}$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights
for episode = 1,  $M$  do
  Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$ 
  for  $t = 1, T$  do
    With probability  $\epsilon$  select a random action  $a_t$ 
    otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$ 
    Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
    Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
    Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$ 
    Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$ 
    Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$ 
    Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  according to equation 3
  end for
end for
  
```



Algorithm 1 DDPG algorithm

```

Randomly initialize critic network  $Q(s, a | \theta^Q)$  and actor  $\mu(s | \theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .
Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$ 
Initialize replay buffer  $R$ 
for episode = 1,  $M$  do
  Initialize a random process  $\mathcal{N}$  for action exploration
  Receive initial observation state  $s_1$ 
  for  $t = 1, T$  do
    Select action  $a_t = \mu(s_t | \theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise
    Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$ 
    Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$ 
    Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$ 
    Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1} | \theta^{\mu'}) | \theta^{Q'})$ 
    Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i | \theta^Q))^2$ 
    Update the actor policy using the sampled policy gradient:
      
$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a | \theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s | \theta^\mu)|_{s_i}$$

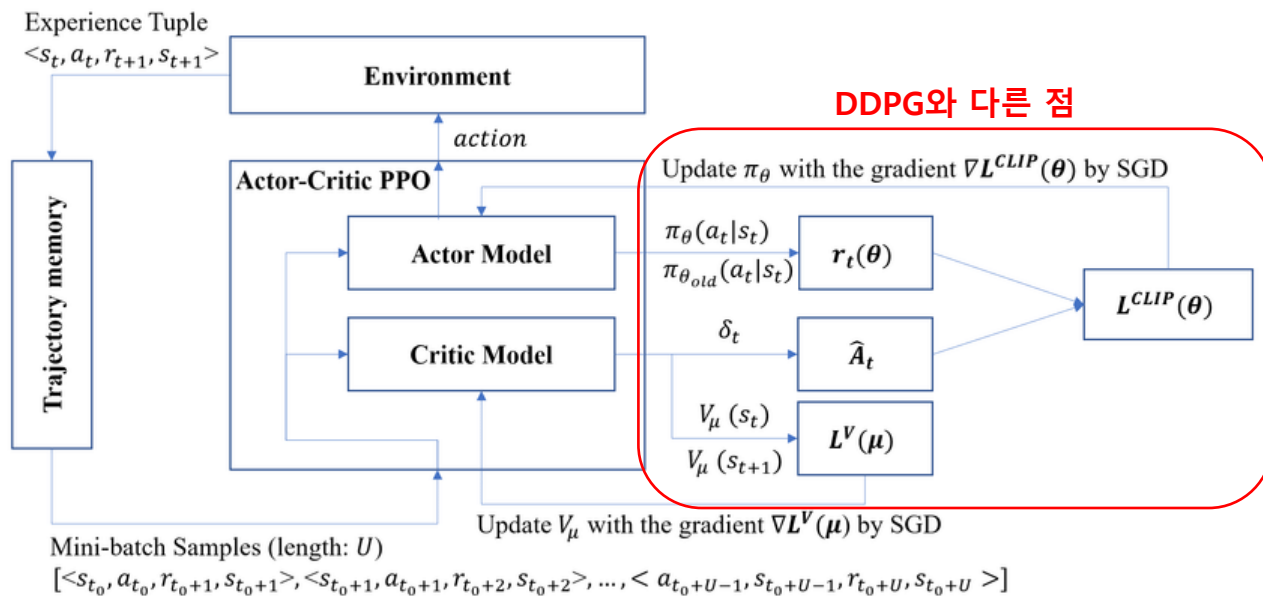
    Update the target networks:
      
$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

      
$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

  end for
end for
  
```

- Discrete -> Continuous action
- Simple CNN structure -> Actor & Critic

PPO (Proximal Policy Optimization)



$$r_t(\theta) = \frac{\pi_\theta(a_t | s_t)}{\pi_{\theta_{old}}(a_t | s_t)}$$

전후 state 확률 비율

$$= L^{CLIP}(\theta) = \mathbb{E}_t \left[\min(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right]$$

$$\hat{A}_t = \delta_t + (\gamma\lambda)\delta_{t+1} + \dots + (\gamma\lambda)^{T-t+1}\delta_{T-1},$$

$$\text{where } \delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$$

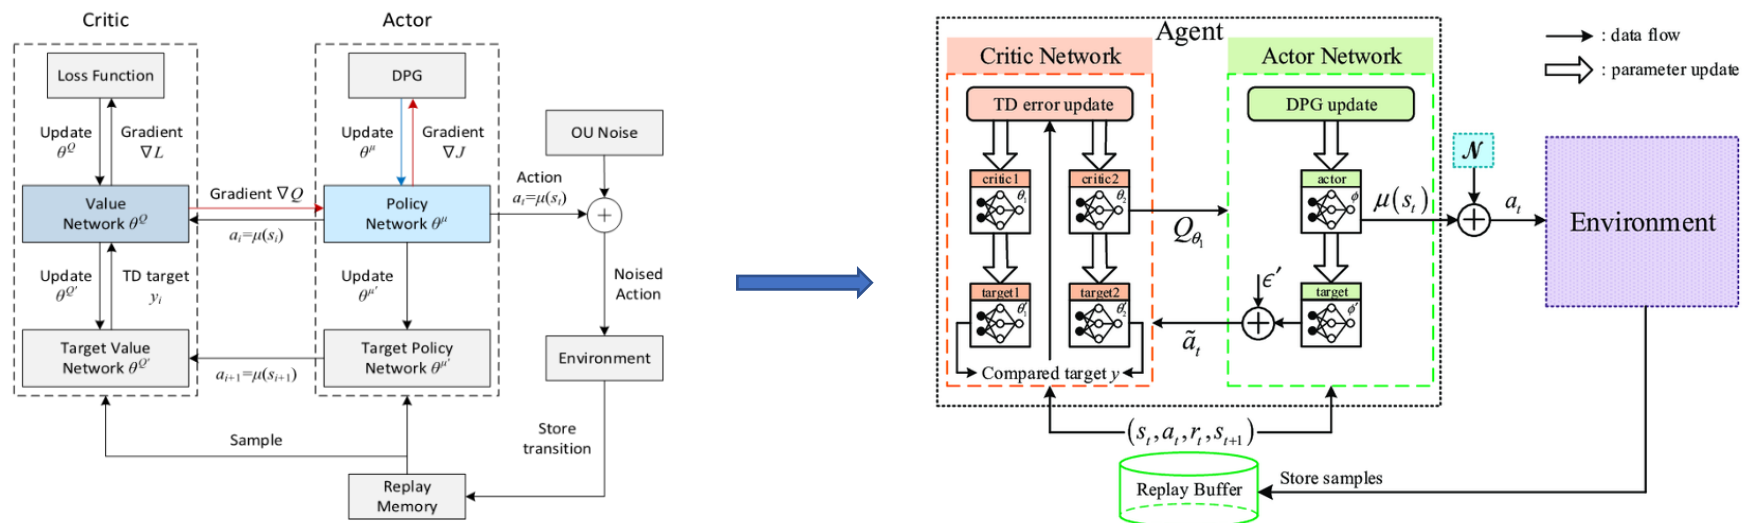
Algorithm 1 PPO, Actor-Critic Style

```

for iteration=1, 2, ... do
  for actor=1, 2, ..., N do
    Run policy  $\pi_{\theta_{old}}$  in environment for  $T$  timesteps
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$ 
  end for
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$ 
   $\theta_{old} \leftarrow \theta$ 
end for
    
```


Recent RL Algorithms

- <DDPG>에서 <TD3>로 발전 (4)



- PPO 변형 알고리즘

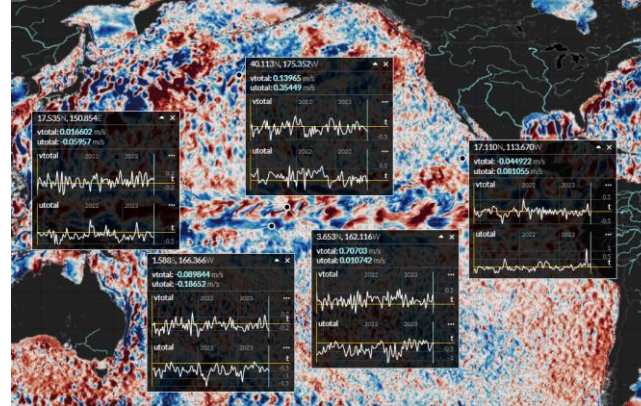
-> Adaptive PPO (학습 과정 중에 학습률 변경)

-> MoPPO (Modified Actor-Critics) (5)

[4] Modified Actor-Critics [Erinc Merdivan, Sten Hanke, Matthieu Geist](https://doi.org/10.48550/arXiv.1907.01298) <https://doi.org/10.48550/arXiv.1907.01298>

[5] Addressing Function Approximation Error in Actor-Critic Methods [Scott Fujimoto, Herke van Hoof, David Meger](https://doi.org/10.48550/arXiv.1802.09477) <https://doi.org/10.48550/arXiv.1802.09477>

Applications



$$\begin{bmatrix} \dot{x} \\ \dot{y} \\ \dot{\psi} \end{bmatrix} = \begin{bmatrix} V \cos \psi + C_x(x, y) \\ V \sin \psi + C_y(x, y) \\ r_c \end{bmatrix}$$

Vessel Motion
(Action)

$$C_{RB}(\nu) = \begin{bmatrix} 0 & 0 & -m(x_g r + \nu) \\ 0 & 0 & \mu \\ m(x_g r + \nu) & -\mu & 0 \end{bmatrix}$$

Drag, Interference Coeff
(State)

Ship Motion Control

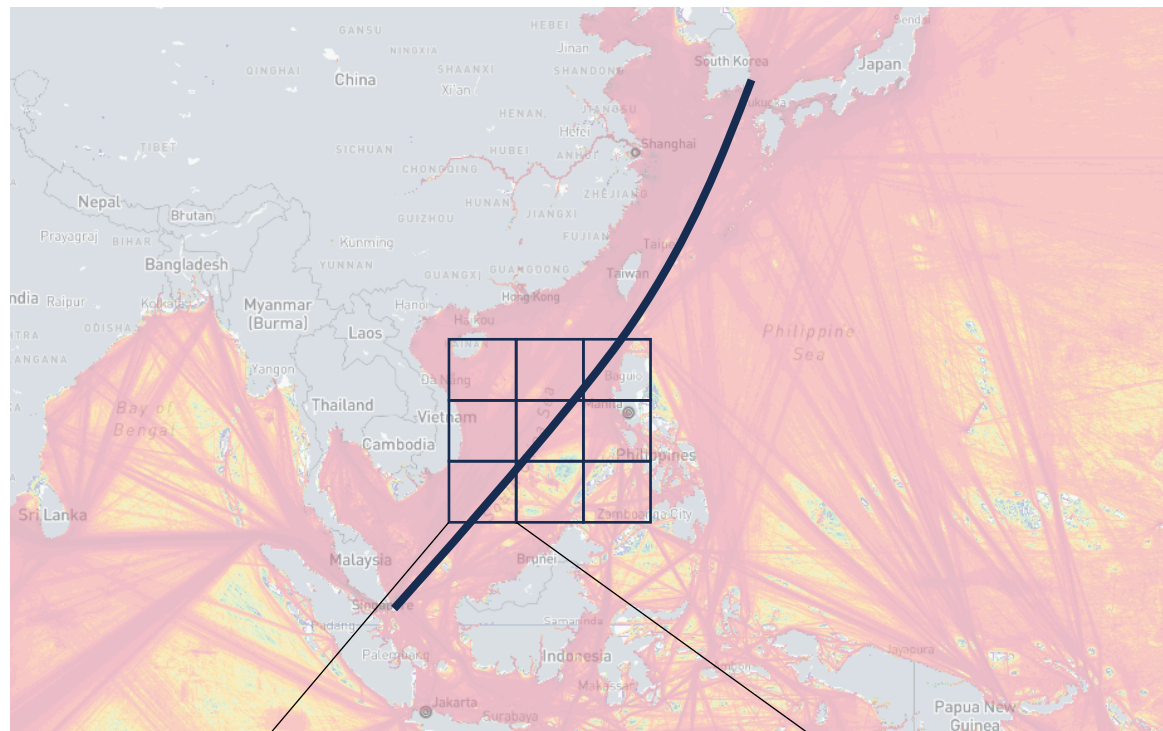
COURSE KEEPING AND ROLL STABILISATION
USING RUDDER AND FINS

Tristan Perez

navigation

our system


path planning



4606 km
(busan -> singapore)

6.1 days
(17knot)

31.484km/h
1knot= 1.852km/h



15 ~ 20 km
each grid

0.5 ~ 1 hour

감사합니다.
